Approximate Matching

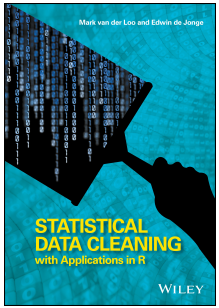

Mark van der Loo and Edwin de Jonge

Statistics Netherlands Research & Development
Twitter: @markvdloo @edwindjonge

ENBES|EESW 2019

Try the code together with your neighbour

`02input/raw_to_input.R`

String distance

Default (Optimal String Alignment distance)

Count number of character deletions, insersions, substitutions and transpositions (of adjacent characters)

```
library(stringdist)
stringdist("Ross Ihaka", "Robert Gentleman")
```

```
## [1] 12
```

Exact Matching with `match`

```
lookup <- c("Alice","Bob","Carol","Danny")
raw    <- c("Bob","Carl","Rob","bob","Dan","Alice")
i <- match(raw, lookup)
data.frame(raw=raw, matched=lookup[i])
```

```
##     raw matched
## 1   Bob     Bob
## 2  Carl    <NA>
## 3   Rob    <NA>
## 4   bob    <NA>
## 5   Dan    <NA>
## 6 Alice   Alice
```

Approximate Matching with `stringdist::amatch`

```
library(stringdist)
j <- amatch(raw, lookup, maxDist=2)
data.frame(raw=raw, matched=lookup[i], amatched=lookup[j])
```

```
##     raw matched amatched
## 1   Bob     Bob       Bob
## 2  Carl    <NA>     Carol
## 3   Rob    <NA>       Bob
## 4   bob    <NA>       Bob
## 5   Dan    <NA>     Danny
## 6 Alice   Alice     Alice
```

$\rightarrow$ Match with closest match, and distance ≤ 2.

Optimal string alignment?

```
stringdist("Robert Gentleman", "Gentleman, Robert")
```

```
## [1] 15
```

```
stringdist("Robert Gentleman", "Ross Ihaka")
```

```
## [1] 12
```

→ OSA wil give a false match (if we allow `maxDist` of 12)

Alternative: cosine distance

```
stringdist("Robert Gentleman","Gentleman, Robert"
           , method="cosine", q=2)
```

```
## [1] 0.1608536
```

```
stringdist("Robert Gentleman","Ross Ihaka"
           , method="cosine", q=2)
```

```
## [1] 0.9139337
```

Notes

- Based on counting co-occurrence of character q-grams (here: pairs).
- Always between 0 and 1

More on `amatch`

```
amatch(x, table, method, maxDist,...)
```

`x`	character data to be matched
`table`	the lookup table with clean values
`method`	string distance type
`maxDist`	Maximum distance allowed (depends on `"method"`!)
`...`	Extra options depending on `"method"`

Example

```
amatch(raw, lookup, method="cosine", maxDist=0.5, q=3)
```

Assignment

Merge data from the `companies` dataset with data from `backbone.csv`.

- Using approximate matching on the `"name"` and `"company"` column.
- Think about and try different distance functions and `maxDist`
- Keep your best solution
- Remove rows that cannot be matched
- Write to `02input/myinput.csv`