



Approximate Matching

Mark van der Loo and Edwin de Jonge

Statistics Netherlands Research & Development

Twitter: @markvdloo @edwindjonge

ENBES|EESW 2019



Try the code together with your neighbour

```
02input/raw_to_input.R
```



String distance

Default (Optimal String Alignment distance)

Count number of character deletions, insertions, substitutions and transpositions (of adjacent characters)

```
library(stringdist)
stringdist("Ross Ihaka", "Robert Gentleman")
```

```
## [1] 12
```



Exact Matching with match

```
lookup <- c("Alice","Bob","Carol","Danny")
raw      <- c("Bob","Carl","Rob","bob","Dan","Alice")
i <- match(raw, lookup)
data.frame(raw=raw, matched=lookup[i])
```

##	raw	matched
## 1	Bob	Bob
## 2	Carl	<NA>
## 3	Rob	<NA>
## 4	bob	<NA>
## 5	Dan	<NA>
## 6	Alice	Alice



Approximate Matching with `stringdist::amatch`

```
library(stringdist)
j <- amatch(raw, lookup, maxDist=2)
data.frame(raw=raw, matched=lookup[i], amatched=lookup[j])
```

##	raw	matched	amatched
## 1	Bob	Bob	Bob
## 2	Carl	<NA>	Carol
## 3	Rob	<NA>	Bob
## 4	bob	<NA>	Bob
## 5	Dan	<NA>	Danny
## 6	Alice	Alice	Alice

→ Match with closest match, and distance ≤ 2 .



Optimal string alignment?

```
stringdist("Robert Gentleman", "Gentleman, Robert")
```

```
## [1] 15
```

```
stringdist("Robert Gentleman", "Ross Ihaka")
```

```
## [1] 12
```

→ OSA wil give a false match (if we allow maxDist of 12)



Alternative: cosine distance

```
stringdist("Robert Gentleman", "Gentleman, Robert",  
           , method="cosine", q=2)
```

```
## [1] 0.1608536
```

```
stringdist("Robert Gentleman", "Ross Ihaka",  
           , method="cosine", q=2)
```

```
## [1] 0.9139337
```

Notes

- Based on counting co-occurrence of character q -grams (here: pairs).
- Always between 0 and 1



More on amatch

```
amatch(x, table, method, maxDist,...)
```

x	character data to be matched
table	the lookup table with clean values
method	string distance type
maxDist	Maximum distance allowed (depends on "method!")
...	Extra options depending on "method"

Example

```
amatch(raw, lookup, method="cosine", maxDist=0.5, q=3)
```



Assignment

Merge data from the `companies` dataset with data from `backbone.csv`.

- Using approximate matching on the "name" and "company" column.
- Think about and try different distance functions and `maxDist`
- Keep your best solution
- Remove rows that cannot be matched
- Write to `02input/myinput.csv`



More on string distances

Main idea

Define a sense of *distance* between two text strings

Distance

A function $d(s, t)$ that takes two arguments and

- returns a nonnegative number,
- returns zero if and only if $s = t$
- is symmetric: $d(s, t) = d(t, s)$
- is the length of a shortest path between s and t : $d(s, t) \leq d(s, u) + d(u, t)$

Note

Some string distances violate one or more of the above assumptions.



Distance types

- Edit based
- q -gram based
- Heuristic



Edit-based string distances

Idea

1. Choose basic steps to alter a string
2. Find the smallest nr of steps that changes s into t
3. The distance equals the nr of steps needed.

Basic steps

- deletion: $hihi \rightarrow hii$
- insertion: $hihi \rightarrow hihh$
- substitution: $hihi \rightarrow hiha$
- transposition: $hihi \rightarrow ihhi$



Edit-based string distances

Distance	Allowed operation			
	substitution	deletion	insertion	transposition
Hamming	✓	✗	✗	✗
LCS	✗	✓	✓	✗
Levenshtein	✓	✓	✓	✗
OSA	✓	✓	✓	✓*
Damerau-Levenshtein	✓	✓	✓	✓

*Substrings may be edited only once.



Example

Levenshtein distance

$$leia \xrightarrow[+1]{sub} lela \xrightarrow[+1]{ins} leela$$

Longest common subsequence distance

$$leia \xrightarrow[+1]{del} lea \xrightarrow[+1]{ins} leea \xrightarrow[+1]{ins} leela$$



q -gram based distances (I)

Algorithm

- Tabulate substrings of length q ($= q$ -gram profile)
- Compute a distance between the profiles

Example

2-gram profile of banana

ba	an	na
1	1	2



q -gram based distances (II)

q -gram distance

Manhattan distance between q -gram profiles

$$\sum_{q\text{-gram}} |n_{q\text{-gram}}(s) - n_{q\text{-gram}}(t)|$$

Cosine-distance

1 minus the cosine of the angle between the profiles

$$1 - \frac{\mathbf{n}(s) \cdot \mathbf{n}(t)}{\|\mathbf{n}(s)\| \|\mathbf{n}(t)\|}$$

Node

- Does not satisfy the 'identity' demand.
- Often one chooses $q = 2$ or $q = 3$



Jaro-Winkler distance

Jaro distance

$$d_j(s, t) = 1 - \frac{1}{3} \left(\frac{m}{|s|} + \frac{m}{|t|} + \frac{m + T}{m} \right)$$

- m number of matching characters (within a window)
- T number of matches that need swapping
- $|s|, |t|$ number of characters in s, t .

Jaro-Winkler distance

$$d_{jw}(s, t) = [1 - p\ell(s, t)]d_j(s, t)$$

- $\ell(s, t)$ length of longest equal prefix (up to 4 characters)
- p a number between 0 and 0.25 (usually 0.1)



Soundex

Algorithm

- Strings are appointed a code: same code means 'sounds the same'
- Equal codes: distance zero, otherwise 1

Example

- Farnsworth → H652
- Fnarswort → H562

Note

- Based on English pronunciation
- Many extensions exist (see the `phonics` R package)



Which one to use?

Considerations

- Fixed versus variable structure/length
- Why would strings differ? (typos, speech, deliberate changes)
- Performance

